What is UTF-8?

UTF-8 is a variable-length character encoding for Unicode. It uses 1 to 4 bytes per character.

It is backward compatible with ASCII, so plain English text produces identical bytes.

UTF-8 is the dominant encoding on the web because it can represent every Unicode character.

Example: the letter A is stored as the single byte 0x41, while the euro sign is stored as three bytes: E2 82 AC.